1 Introduction

Recurrent models typically factor computation along the symbol positions of the input and output sequences. We propose the Transformer, a model architecture eschewing recurrence.

6 Results

On the WMT 2014 English-to-French translation task, our big model achieves a BLEU score of 41.0, outperforming all of the previously published single models.